

Deploying Physical Intelligence $\pi_{0.5}$ on the Unified Engine

Apex Compute
hunlu@apexcompute.com

1 The $\pi_{0.5}$ Model

$\pi_{0.5}$ is a vision–language–action (VLA) policy of roughly three billion parameters that maps camera observations and a language instruction directly to robot actions. Each inference emits a 10×7 action chunk: ten future control steps, each a 7-dimensional command giving the robot’s degrees of freedom (arm motion plus gripper). The computation splits into three stages run in sequence—a vision encoder, a language prefix, and a flow-matching action expert.

1.1 Vision encoder

Each image is encoded independently by a SigLIP ViT; $\pi_{0.5}$ takes up to three slots (scene, wrist, one padded). Every 224×224 image becomes 256 patch tokens of width 1152, run through $L_v=27$ bidirectional layers (16 heads, head-dim 72, MLP width 4304). Re-run per image and never cached, it costs ≈ 656 GFLOP.

1.2 Language prefix

The patch tokens are projected to $d=2048$, concatenated with the instruction tokens, and processed once by a Gemma-2B backbone: $L_p=18$ layers of grouped-query attention (8 query, 1 KV head, head-dim 256), RMSNorm, RoPE, and a gated SiLU MLP ($I=16384$). At prefill length $M_p=832$ one layer costs

$$2M_p d d_q + 4M_p d d_{kv} + 2M_p d_q d + 4M_p^2 d_q + 6M_p d I, \quad (1)$$

($d_q=2048$, $d_{kv}=256$), dominated by the MLP; over 18 layers the prefix is ≈ 3400 GFLOP (82% of the model). Each layer’s K/V is cached and reused by every denoise step.

1.3 Action expert

A second, smaller Gemma (~ 300 M; $H=1024$, MLP 4096, 18 layers) turns noise into the action chunk by flow matching: from Gaussian noise it integrates a learned velocity field v_θ over $N=10$ Euler steps ($\Delta t = -0.1$),

$$x_{t+\Delta t} = x_t + \Delta t v_\theta(x_t, t, \text{KV}_{\text{prefix}}), \quad (2)$$

each step attending over the cached prefix K/V and reading out actions through a $(1024 \rightarrow 32)$ head. At only ≈ 75 GFLOP (1.8%) it is arithmetically small but the sole sequential stage, and its per-step error accumulates through the integrator.

Table 1: Per-stage parameter and FLOP budget (per inference). *Effective* uses true model dims; *FPGA* is the padded shape the transformer engine actually issues. The 832-token prefix is 64-aligned (no padding); the vision tower pads its attention head-dim (72→128) and MLP (4304→4352), ~30% extra issued work; the action expert pads its query horizon 10→64 and action width 32→64, inflating its work ~6.5× — so overall the FPGA issues ~15% more FLOPs than the model nominally needs.

Stage	Params	Eff. GFLOP	FPGA GFLOP
Vision tower (SigLIP, 3 slots)	~0.41 B	655.9	850.7
Language prefix (Gemma-2B)	~2.0 B	3399.8	3399.8
Action expert (Gemma-300M)	~0.31 B	74.7	483.5
Total	~2.7 B	4130.4	4734.0

2 Porting

The Unified Engine v1 is a soft accelerator IP block whose compute core is the Apex Compute Transformer Engine; it can be deployed on any sufficiently large FPGA and operates on 64-element alignment. In this work a single transformer engine is targeted to a Xilinx Kintex-7 with 4 GB of attached DRAM. These two constraints—64-alignment and the 4 GB DRAM budget—drive the porting transformations below.

IF4 weights / DRAM budget. The ~13 GB full-precision checkpoint does not fit the 4 GB DRAM budget. All matmul weights use 4-bit block quantization (block 64, weight-only): every 64-element block stores 64 INT4/FP4 codes plus *one* bf16 scale, so a quantized matrix costs 4.25 bits/weight (4 + 16/64, i.e. 3.76× smaller than bf16 rather than 4×), while norm gains and biases stay bf16 outright. That shrinks the three stacks to ~1.56 GB. Weights are dequantized to bf16 immediately before each multiply–accumulate, so IF4 is a *footprint enabler*, not a *speedup* — compute is bf16 throughout.

Vision head-dim pad. Head-dim 72 is not 64-/128-aligned, so the attention path runs at a padded head-dim of 128 — weights pre-padded 72→128 across the $Q/K/V/O$ projections and the attention matmuls — and is contracted back via a precomputed 128×72 selection matrix; the vision MLP pads 4304→4352. This padding is real issued work, ~30% over the effective FLOPs (Table 1).

Prefix masking. When an image slot is unused it can be dropped from the prefill, compressing 832→576 (bit-identical, ~31% less prefill work). Rotary positions use cumulative-sum-of-mask indices, not **arange**, so dropping a slot does not shift phase. The non-causal/causal mixed attention mask is supplied as a static bias tensor in DRAM. The measured runs in Sec. 4 keep the full 832-token prefill so the 3-slot path is exercised end-to-end (Sec. 3.2).

Denoise loop. All 10×18 steps are unrolled into one program issued with a single execute call; AdaRMSNorm conditioning is a matmul to $3H$ split into scale/shift/gate.

3 Quantization

3.1 The Quantization Floor

IF4 is the accelerator’s 4-bit weight format: a per-64-block hybrid of signed INT4 (uniform grid $[-8, 7]$) and FP4 E2M1 (non-uniform codebook $\{0, \pm 0.5, \pm 1, \pm 1.5, \pm 2, \pm 3, \pm 4, \pm 6\}$, finer near zero).

At pack time each block is quantized *both* ways and the lower-MSE variant kept; the choice rides in the *sign* of the block’s scale (neg $\rightarrow$ INT4, pos $\rightarrow$ FP4), so the on-chip dequantizer picks the decode path from one bit, inline before the MAC—no metadata. INT4 suits dense blocks, FP4 the peaky/outlier ones. (Results below use the INT4 variant, q4_64.)

We measure the accuracy cost against a bf16 reference (identical torch run, same inputs/noise, diffed per stage; Table 2, $\text{SNR}=20\log_{10}(\|b\|/\|a-b\|)$, masked rows excluded, sample 0)—the theoretical quantization floor, isolated from any hardware effect. Vision is nearly untouched (25.4 dB); the prefix caches degrade with depth, V worse than K (10.7 vs. 16.0 dB mean); and while each denoise step’s velocity holds up on its own (23.5 dB), the Euler loop compounds that per-step error through Eq. (2)— x_t walks 45 $\rightarrow$ 16.3 dB (Table 3), concentrating in low-magnitude action dims (dim 1, 1.5 dB) while task-critical dims hold (dim 0 24.9, gripper 40.7 dB). SNR is a diagnostic; the acceptance gate is the closed-loop LIBERO agreement (Sec. 3.2).

Table 2: IF4 vs. bf16 per section (torch reference, sample 0).

Model section	SNR (dB)	Range (dB)
Vision encoder (3 cameras)	25.37	24.3–26.6
Prefix K-cache (18 layers)	15.96	12.5–19.6
Prefix V-cache (18 layers)	10.65	7.0–15.4
Denoise velocity v_t (10 steps)	23.49	11.5–26.2
End-to-end action (10$\times$7)	16.31	MSE 4.05×10^{-3}

Table 3: Final-action detail: per-dimension SNR (left) and error accumulation of the running state x_t across Euler steps (right). Sample 0.

Action dim	SNR (dB)	Euler step	x_t SNR (dB)
dim 0 (primary)	24.91	step 0	44.95
dim 1	1.50	step 3	30.13
dim 2	10.08	step 5	24.15
dim 3	6.73	step 7	20.53
dim 4	4.65	step 8	19.26
dim 5	17.61	step 9	16.31
dim 6 (gripper)	40.73	(final)	16.31

3.2 LIBERO

LIBERO is a standard tabletop manipulation benchmark for evaluating language-conditioned robot policies: a suite of pick-and-place and articulated tasks in simulation, scored by closed-loop task success over many episodes. Its observation uses **two** camera views—a fixed scene camera and a wrist-mounted camera. $\pi_{0.5}$, however, exposes **three** image slots, and we process all three: the two real LIBERO cameras plus a third all-zero (attention-masked) slot. Encoding the empty slot is intentional—other embodiments and benchmarks drive the policy with a genuine third camera for their robot actions, so keeping the full 3-slot path exercised on hardware makes the port drop-in for those settings at no correctness cost (the masked slot contributes nothing to the output).

An episode is many inferences. Each model call in Sec. 1 predicts a 10 $\times$ 7 chunk—ten future control steps. LIBERO runs *closed-loop*: the simulator executes the first $k=10$ steps of that chunk, then feeds the resulting new observation back to the policy for a fresh chunk, and repeats. So a

single task *episode* is a *sequence* of inferences—one every k sim steps—not a one-shot prediction, and the episode runs until the task’s goal predicate fires (success) or a per-suite step cap is hit (failure). An episode therefore takes $\lceil \text{steps}/k \rceil$ inferences: the successful runs in Table 4 are ~ 10 –16 inferences, the capped failures ~ 23 –29. This is what makes the closed-loop test meaningful and also much heavier than a single forward pass—every one of those inferences runs the full vision + prefix + 10-step denoise pipeline on hardware.

Per-tensor SNR is only a diagnostic; the real acceptance test is whether the quantized policy *acts* the same across that whole loop. Table 4 is the closed-loop comparison—the FPGA IF4 policy vs. the CUDA bf16 reference over 10 prompts (9 `libero_spatial` + 1 `libero_object`), one episode each, with identical seed and fixed denoise noise so both platforms see byte-identical observations. We score each by task success and by the per-episode step count.

Table 4: LIBERO closed-loop agreement: FPGA (IF4) vs. CUDA (bf16), 10 prompts, 1 episode each, fixed seed/noise. $\checkmark$ =success, $\times$ =failure; the number is episode length in sim steps. Δ is FPGA – CUDA steps on successes (failures run to the per-suite step cap on both platforms, so Δ is not meaningful there). The two policies reach the *same* outcome on every prompt.

Prompt	FPGA (IF4)	CUDA (bf16)	Δ steps	Agree
0	$\checkmark$ 94	$\checkmark$ 95	−1	$\checkmark$
1	$\checkmark$ 160	$\checkmark$ 126	+34	$\checkmark$
2	$\checkmark$ 123	$\checkmark$ 115	+8	$\checkmark$
3	$\checkmark$ 108	$\checkmark$ 109	−1	$\checkmark$
4	$\times$ 230	$\times$ 230	—	$\checkmark$
5	$\times$ 290	$\times$ 290	—	$\checkmark$
6	$\checkmark$ 121	$\checkmark$ 120	+1	$\checkmark$
7	$\times$ 230	$\times$ 230	—	$\checkmark$
8	$\checkmark$ 107	$\checkmark$ 106	+1	$\checkmark$
9	$\times$ 230	$\times$ 230	—	$\checkmark$
Rate	6/10	6/10		10/10

Takeaway. IF4 is not just numerically close—it is behaviorally identical. Across all 10 prompts the FPGA IF4 policy and the bf16 reference reach the *same* success/failure outcome (10/10 agreement, 6/6 shared successes), and on the successful episodes their trajectory lengths differ by only a handful of steps (Δ within a few percent, +34 on one longer reach). The 16.3 dB end-to-end SNR floor of Sec. 3 sounds alarming in isolation, but the low-SNR components are small-magnitude action dimensions that do not change what the robot *does*—the task-critical axes and gripper are preserved, and the closed-loop behavior confirms it. Put as a benchmark number: quantization *did not regress task performance at all* — same 6/10 success rate, same outcome on every prompt, successful trajectories within a few control steps — so the $\sim 8\times$ checkpoint compression is free at the task level on this suite.

4 Throughput: FPGA vs. CUDA

We time each stage on both platforms and report it in the same form: per-section wall-clock and the achieved rate (section FLOPs $\div$ time). FPGA times come from the engine’s own instrumentation. The GPU baseline is openpi’s JAX/XLA Pi0 — the deployment path, sustaining 130.2 ms/inference (7.68 Hz), batch 1 — broken into the same three stages by jitting each separately and forcing device completion (`block_until_ready`) at the boundaries. Because these stages are dispatched separately

rather than fused into the deployed `lax.while_loop`, their sum (135.4ms) slightly exceeds the fused end-to-end (130.2ms); the per-stage split is used for the breakdown, the fused number for the headline.

Table 5: RTX 5070 (bf16, openpi JAX/XLA) per-section timing and achieved throughput. Batch 1, 10 denoise steps, sample 0. Per-stage times are separately-jitted; End-to-end is the fused `policy.infer`.

Model section	Time (ms)	GFLOP/s
Vision encoder	19.41	33 793
Prefix (Gemma-2B)	79.29	42 877
Action expert (denoise)	36.70	2 035
End-to-end	130.15	31 736
Throughput (inf/s)	7.68	

Table 6: FPGA (IF4) per-section timing and achieved throughput (Unified Engine v1 on Xilinx Kintex-7). Batch 1, 10 denoise steps, sample 0, 3-slot vision, 832-token prefix. GFLOP/s and % peak are on the hardware-issued (padded) FLOPs; theoretical peak is 26.6 GFLOP/s (one Apex Compute Transformer Engine at 208 MHz).

Model section	Time (s)	Share	Eff. GFLOP	HW GFLOP	GFLOP/s (hw)	% peak
Vision encoder	35.4	18.3%	655.9	850.7	24.0	90%
Prefix (Gemma-2B)	137.6	71.1%	3399.8	3399.8	24.7	93%
Action expert (denoise)	20.6	10.6%	74.7	483.5	23.5	88%
End-to-end	193.6	100%	4130.4	4734.0	24.5	92%
Throughput (inf/s)				0.0052		

Table 7: Silicon (IF4, *estimated*) per-section timing and achieved throughput — projected ASIC implementation of the Unified Engine v1 IP at 1.6 GHz with 100 GB/s DRAM, scaling the Apex Compute Transformer Engine count in two chip variations — **Chip Variation 1** (8 engines, ~5 W) and **Chip Variation 2** (16 engines, wattage TBD) — 1638.4 / 3276.8 GFLOP/s aggregate peak. Batch 1, 10 denoise steps. Linear $61.5\times$ / $123\times$ compute scaling of the measured FPGA run at its achieved %-of-peak; neither variation is bandwidth-bound (see below).

Model section	Chip Variation 1 (8 engines, ~5 W)		Chip Variation 2 (16 engines, wattage TBD)	
	Time (ms)	GFLOP/s	Time (ms)	GFLOP/s
Vision encoder	575	1 479	288	2 958
Prefix (Gemma-2B)	2 236	1 520	1 118	3 041
Action expert (denoise)	335	1 444	167	2 889
End-to-end	3 146	1 505	1 573	3 010
Throughput (inf/s)		0.32		0.64

Why the projection is linear. Going from one Apex Compute Transformer Engine at 208 MHz to 8 engines at 1.6 GHz multiplies compute by $61.5\times$ ($26.6 \rightarrow 1638.4$ GFLOP/s peak) but DRAM bandwidth by only $16.7\times$ ($\sim 6 \rightarrow 100$ GB/s), so each section survives as compute-bound only if its arithmetic intensity clears the ASIC ridge point of $1638.4/100 \approx 16$ FLOP/byte. Traffic here is dominated by streaming the IF4 weights (4.25 bits/weight, Sec. 2): the prefix reads

~ 1.06 GB once for its 3399.8 GFLOP (~ 3200 FLOP/byte), the vision tower ~ 0.65 GB over three slots (~ 1300 FLOP/byte), and the worst case — the action expert, whose ~ 0.16 GB are streamed on *every* denoise step against just ~ 48 GFLOP of padded work per step — still lands at $\gtrsim 200$ FLOP/byte with activation and KV-cache traffic included. All three sections sit an order of magnitude above the ridge (worst-section demand < 10 GB/s of the 100 available), and doubling to 16 engines (Chip Variation 2) only lifts the ridge to ≈ 33 FLOP/byte — still $\geq 6\times$ clear. So neither variation is bandwidth-bound, and times extrapolate linearly at the FPGA’s measured 88–93% of peak, assuming ideal engine scaling: 3.15 s per inference (0.32 inf/s) for Chip Variation 1 and 1.57 s (0.64 inf/s) for Chip Variation 2 — $61.5\times$ / $123\times$ the FPGA, within $24\times$ / $12\times$ of the RTX 5070’s 7.68 inf/s. Either way the engine executes the whole vision–prefix–denoise pipeline autonomously: host work per action chunk is an embedding gather on the way in and a 10×7 readback on the way out, so the platform’s CPU stays free for the robot stack.

Power. The RTX 5070 is a 250 W-TGP board; Chip Variation 1 is specified at ~ 5 W, and Chip Variation 2’s wattage is TBD. Table 8 collects the platform comparison.

Table 8: Platform comparison: measured RTX 5070 and Unified Engine v1 FPGA runs, and the two projected chip variations. Achieved GFLOP/s counts effective FLOPs on the GPU and hardware-issued FLOPs on the Unified Engine platforms (Tables 5–7); peaks are bf16 rates.

	RTX 5070	FPGA (Kintex-7)	Chip Var. 1	Chip Var. 2
Compute	6144 CUDA cores	1 engine	8 engines	16 engines
Clock	~ 2.5 GHz	208 MHz	1.6 GHz	1.6 GHz
Peak GFLOP/s (bf16)	61 740	26.6	1 638	3 277
DRAM bandwidth	672 GB/s	~ 6 GB/s	100 GB/s	100 GB/s
Power	250 W (TGP)		~ 5 W	TBD
Achieved GFLOP/s	31 736	24.5	1 505	3 010
Time / inference	0.130 s	193.6 s	3.15 s	1.57 s
Throughput	7.68 inf/s	0.0052 inf/s	0.32 inf/s	0.64 inf/s

Config. FPGA: Unified Engine v1 IP (one Apex Compute Transformer Engine) on Xilinx Kintex-7 at 208 MHz (26.6 GFLOP/s theoretical peak), IF4 weights / bf16 compute, 4 GB DRAM at ~ 6 GB/s. GPU: RTX 5070 (6144 CUDA cores, 250 W TGP), bf16, 12 GB GDDR7 at 672 GB/s, openpi JAX/XLA. Silicon: Chip Variation 1 — 8 engines at 1.6 GHz, ~ 5 W; Chip Variation 2 — 16 engines at 1.6 GHz, wattage TBD; both 100 GB/s DRAM. All: batch 1, 10 denoise steps, sample 0.

The entire $\pi_{0.5}$ VLA — vision, Gemma-2B prefix, and the 10-step flow-matching expert — runs on the Unified Engine v1. The recurring theme is the split between where FLOPs concentrate (the compute-bound prefix) and where latency does (the sequential, FLOP-light denoise loop), and how IF4 trades a controlled, section-localized accuracy cost — one with no measured closed-loop task regression (Sec. 3.2) — for fitting a ~ 13 GB model in a sub-4 GB budget. Tables 1–8 complete the characterization.